

# Dynamic trees as search trees via Euler tours, applied to the network simplex algorithm

Robert E. Tarjan<sup>1</sup>

*Department of Computer Science, Princeton University, Princeton, NJ 08544 and NEC Research Institute,  
4 Independence Way, Princeton, NJ 08540, USA*

Received 15 September 1995; revised manuscript received 17 May 1996

---

## Abstract

The *dynamic tree* is an abstract data type that allows the maintenance of a collection of trees subject to joining by adding edges (*linking*) and splitting by deleting edges (*cutting*), while at the same time allowing reporting of certain combinations of vertex or edge values. For many applications of dynamic trees, values must be combined along paths. For other applications, values must be combined over entire trees. For the latter situation, an idea used originally in parallel graph algorithms, to represent trees by Euler tours, leads to a simple implementation with a time of  $O(\log n)$  per tree operation, where  $n$  is the number of tree vertices. We apply this representation to the implementation of two versions of the network simplex algorithm, resulting in a time of  $O(\log n)$  per pivot, where  $n$  is the number of vertices in the problem network. © 1997 The Mathematical Programming Society, Inc. Published by Elsevier Science B.V.

---

## 1. Introduction

Consider a collection of unrooted trees, each initially a single vertex and no edges, on which two structural update operations are allowed:

*link*( $\{v, w\}$ ): Combine the trees containing vertices  $v$  and  $w$  by adding the edge  $\{v, w\}$ . This operation does nothing if  $v$  and  $w$  are already in the same tree.

*cut*( $\{v, w\}$ ): Break the tree containing edge  $\{v, w\}$  in two by deleting the edge  $\{v, w\}$ . This operation does nothing if  $\{v, w\}$  is not an existing tree edge.

---

<sup>1</sup> Research at Princeton University partially supported by the National Science Foundation, Grant No. CCR-8920505, and the Office of Naval Research, Contract No. N0014-91-J-1463. Work during a visit to M.I.T. partially supported by ARPA Contract No. 14-95-1-1246. Email: ret@cs.princeton.edu.

Suppose further that each tree vertex  $v$  has an associated real value, denoted by  $val(v)$ , and that the following operations on values are allowed:

*find-val*( $v$ ): return  $val(v)$ .

*find-min-val*( $v$ ): return a vertex of minimum value in the tree containing vertex  $v$ .

*change-val*( $v, x$ ): set  $val(v)$  equal to  $x$ .

*add-val*( $v, x$ ): add  $x$  to  $val(w)$  for each vertex  $w$  in the tree containing  $v$ .

Section 2 of this paper presents a simple implementation of these six tree operations with a running time of  $O(\log n)$  per operation, where  $n$  is the number of vertices in the tree or trees involved in the operation. The idea used is to linearize each tree by constructing an *Euler tour* that traverses each edge once in each direction and includes one stop at each vertex, and to represent such a tour by a search tree. Linking and cutting of trees translate into simple combinations of catenation and splitting operations on tours, and on the corresponding search trees. The  $O(\log n)$  time bound per tree operation is amortized if splay trees [22] are used in the representation and worst-case if balanced search trees are used. Section 3 describes the use of the Euler tour structure in the implementation of two versions of the network simplex algorithm.

The Euler Tour representation of trees was originally used in fast parallel graph algorithms [26]. Later, Miltersen et al. [18] adapted it to represent trees subject to linking and cutting but without vertex values. Henzinger and King [16] independently adapted it to represent trees each of whose vertices has an associated list and an associated value which is the size of the associated list; values are combined using sum instead of minimum. Both of these representations use balanced search trees.

Additional related work deals with a class of dynamic trees in which vertex or edge values are combined along paths, rather than over an entire tree. Dynamic trees of this kind arise in various network flow algorithms [10–13,21,23,25] and in other settings [4,6]. Two different representations of such trees have been proposed. The first, by Sleator and Tarjan [21,22,25] decomposes each tree into vertex-disjoint paths and represents these paths by search trees, either biased search trees [3] or splay trees [22]. With the former representation the time per tree operation can be made  $O(\log n)$  in the worst case; with the latter representation the time per tree operation is  $O(\log n)$  amortized over a worst-case sequence of operations. Frederickson [7,8] proposed a rather different representation, called the *topology tree*, which is related to the rake and compress operations used in parallel tree processing [17]. This representation, too, has an  $O(\log n)$  worst-case time bound per tree operation.

Both the Sleator–Tarjan representation and the Frederickson representation can be applied to the problem we consider here, achieving the same  $O(\log n)$  time bound (see e.g. [10]). But we regard these solutions as inferior, for two reasons. First, both data structures must be extended to handle tree vertices of arbitrary degree. Second, both structures, even without the unbounded degree extension, are noticeably more complicated than the Euler tour structure, which ultimately is just a straightforward application of search trees, and is likely to be easier to implement and more efficient in practice.

## 2. Trees as tours and tours as search trees

The representation we describe in this section is a small variant of the one proposed by Mittersen et al. [18] and independently by Henzinger and King [16] for two slightly simpler applications. To represent a dynamic tree  $T$ , we replace each edge  $\{v, w\}$  of  $T$  by two arcs (directed edges)  $(v, w)$  and  $(w, v)$ , and add a loop  $(v, v)$  for each vertex  $v$ . The result is a directed graph such that each vertex has in-degree equal to its out-degree. Such a graph has at least one, and in general many, *Euler tours*: cycles that contain each arc exactly once. We represent the tree by one such tour, broken at an arbitrary place to make it into a list of the arcs.

With this representation, linking and cutting each translate into a fixed set of catenation and splitting operations on lists. Specifically, suppose we wish to perform *link*  $(\{v, w\})$ . Let  $T_1$  and  $T_2$  be the trees containing  $v$  and  $w$  respectively, and let  $L_1$  and  $L_2$  be the lists representing  $T_1$  and  $T_2$ . We split  $L_1$  just after  $(v, v)$ , into lists  $L_1^1$ ,  $L_1^2$ , and we split  $L_2$  just after  $(w, w)$ , into  $L_2^1$ ,  $L_2^2$ . Then we form the list representing the combined tree by catenating the six lists  $L_1^2$ ,  $L_1^1$ ,  $[(v, w)]$ ,  $L_2^2$ ,  $L_2^1$ ,  $[(w, v)]$  in order. Thus linking takes two splits and five catenations; two of the latter are the special case of catenation with singleton lists.

Similarly, suppose we wish to perform *cut*  $(\{v, w\})$ . Let  $T$  be the tree containing  $\{v, w\}$ , represented by list  $L$ . We split  $L$  before and after  $(v, w)$  and  $(w, v)$ , into  $L^1$ ,  $[(v, w)]$ ,  $L^2$ ,  $[(w, v)]$ ,  $L^3$  (or symmetrically  $L^1$ ,  $[(w, v)]$ ,  $L^2$ ,  $[(v, w)]$ ,  $L^3$ ). The lists representing the two trees formed by the cut are  $L^2$  and the list formed by catenating  $L^1$  and  $L^3$ . Thus cutting takes four splits (of which two are the special case of splitting off one element) and one catenation.

By using a simple doubly-linked list representation, we can implement *link* and *cut* to take constant time; indeed, if we make the lists circular, we can save a few pointer updates. But this begs the question of how to handle vertex values. With the representation just presented, both *find-min-val* and *add-val* take time proportional to the tree size, since each requires a scan of the entire corresponding list.

We avoid this inefficiency by representing each list as a search tree. For definiteness, we shall describe a solution based on splay trees [22], a form of self-adjusting binary search tree, although any kind of search tree, such as red–black trees [15,25], AVL trees [1], or B-trees [2], will suffice. Since the required operations on search trees are well-known and straightforward, we shall be very sketchy in the presentation. Detailed discussions of splay trees can be found in [22,25].

Each arc in an Euler tour list becomes a node in the splay tree representing the list. Linear order in the list corresponds to symmetric order in the tree. Each arc has an associated value. A loop  $(v, v)$  has value  $val(v)$ ; a nonloop  $(v, w)$  has value infinity. To handle *add-val* operations efficiently, we store values implicitly, in difference form. Specifically, a node  $x$  of a splay tree has stored with it

$$\text{dif-val}(x) = \begin{cases} \text{val}(x) & \text{if } x \text{ is a splay tree root,} \\ \text{val}(x) - \text{val}(p(x)) & \text{if } x \text{ is a nonroot, where } p(x) \text{ is the parent of } x. \end{cases}$$

In order to keep track of tree minima efficiently, we need to store one additional value at each splay tree node. For such a node  $x$ , we define  $\text{min-val}(x)$  to be  $\min\{\text{val}(y) \mid y \text{ is a descendant of } x \text{ in its splay tree}\}$ . What we actually store with  $x$  is  $\text{dif-min-val}(x) = \text{val}(x) - \text{min-val}(x)$ .

Implementation of the various dynamic tree operations relies on *splaying*, which is the fundamental restructuring operation on splay trees. Splaying moves a designated node to the root of the splay tree by performing a sequence of local restructurings called *rotations*. The amortized time for a splay operation on an  $n$ -node tree is  $O(\log n)$ , including the time to update *dif-val* and *dif-min-val*. Each of the four dynamic tree operations *find-val*, *find-min-val*, *change-val*, and *add-val* takes one splay plus a constant amount of additional work; in the case of *find-min-val* the path along which the splaying takes place is found by a search that is guided by following zero values of *dif-min-val* (see [22]). A list catenation or splitting operation also requires a single splay plus a constant amount of additional work. The implementation details of these operations can be found in [22,25]. Thus all six dynamic tree operations can be performed in  $O(\log n)$  amortized time.

Thus in two steps we have obtained a representation of dynamic trees as search trees with an amortized time bound of  $O(\log n)$  per dynamic tree operation. If we use balanced search trees in place of splay trees, the  $O(\log n)$  time bound becomes worst-case instead of amortized.

### 3. Implementation of network simplex algorithms

We use the Euler tour data structure to implement two versions of the primal network simplex algorithm. The first use is in the Goldfarb–Hao algorithm for the maximum flow problem [14]. The second is in Orlin’s new strongly polynomial algorithm for the minimum-cost flow problem [19,20]. In each case, we obtain an amortized time bound of  $O(\log n)$  per pivot. Both algorithms require two uses of dynamic trees. One is to maintain the residual capacities of arcs. For this purpose a variant of one of the Sleator–Tarjan dynamic tree implementations [21,22] that allows two values per arc suffices. The required changes to the Sleator–Tarjan structure are specified in detail in [23], as is the use of this structure in the network simplex algorithm applied to the minimum-cost flow problem. The use of this structure in the Goldfarb–Hao algorithm is described by Goldberg, et al. [10].

The second need for dynamic trees is to maintain what are in effect dual variables; specifically, shortest path lengths in the case of the Goldfarb–Hao algorithm and “potentials” in the case of the Orlin algorithm. We shall describe how to use the Euler tour structure (in place of a modified Sleator–Tarjan structure) to maintain these values.

We assume familiarity with the relevant parts of [10,14,19,20]; the reader seeking a complete understanding of the network simplex algorithm in these two settings should consult these works.

We base our discussion of the Goldfarb–Hao algorithm on the presentation in [10]. The algorithm maintains a pair of trees  $S, Z$ , with the source vertex  $s$  in  $S$  and the sink vertex  $t$  in  $Z$  ([10], Section 3). Each vertex  $v$  has an associated label, which is a shortest path distance, specifically “the minimum number of pseudoresidual arcs on a path of pseudoresidual arcs from  $s$  to  $v$ ” ([10], Section 4). Each pivot (the elementary step of the algorithm) requires performing the following operations on  $S$ ,  $Z$ , and the vertex labels:

1. Find a vertex  $w$  of minimum label in tree  $Z$ .
2. Link  $S$  and  $Z$  into a tree  $T$  by adding an edge  $\{v, w\}$ .
3. Cut  $T$  into a new pair  $S, Z$  by deleting some edge  $\{x, y\}$ .
4. Update the labels of certain vertices.

We use the data structure of Section 2, with the value of each vertex being its label. Step 1 takes a *find-min-val* operation, step 2 a *link*, step 3 a *cut*, and step 4 one *change-val* per updated label. The total number of label updates over the entire algorithm is  $O(nm)$ , as is the number of pivots [10,14], where  $n$  and  $m$  are the numbers of vertices and edges in the problem network, respectively. It follows that the amortized time per pivot is  $O(\log n)$ , and the total running time of the algorithm is  $O(nm \log n)$ . This matches the bound of Goldberg, et al. for the Goldfarb–Hao algorithm but considerably simplifies one part of the implementation. In this application the *add-val* operation is unnecessary, which allows us to simplify the representation of Section 2 if we wish by storing vertex values explicitly rather than in difference form.

The use of the Euler tour structure in Orlin’s algorithm [19,20] is more complicated. In particular, we must maintain four separate values for each vertex, instead of just one, and we need one additional data structure, a heap for each vertex. Orlin’s algorithm maintains a *basis tree*  $T$  that spans the vertices of the problem network. Tree  $T$  has a designated root vertex  $v$ ; all edges of  $T$  are regarded as being directed toward the root ([19,20], Section 2). The algorithm maintains a real-valued *potential*  $\pi_i$  for each vertex  $i$ . Each arc  $(i, j)$  in the network, including the arcs of  $T$  and their reversals, has a *cost*  $c_{ij}$  that is part of the problem specification and a *reduced cost*  $c_{ij}^\pi = c_{ij} - \pi_i + \pi_j$  ([19,20], Section 1). (This definition of reduced cost uses the negative of the potential as compared to some other such definitions; see e.g. [12,13,25].) The algorithm maintains the property that for each tree arc  $(i, j)$  (directed toward the root),  $c_{ij}^\pi \leq 0$ . ([19,20], Section 3).

The algorithm consists of a sequence of *phases*. It maintains a *scale factor*  $\varepsilon$  that remains constant during a phase and decreases between phases. We shall discuss the performance of a single phase of the algorithm.

We need some terminology. Let us call an arc  $(i, j)$  of  $T$  a *zero arc* if  $c_{ij}^\pi = 0$ . Our implementation maintains a collection of trees, called *zero trees*, the edges of which correspond to some of the zero arcs of  $T$ . (Not all zero arcs of  $T$  need be represented in the zero trees, because changes in vertex potentials can cause arcs to become zero arcs

without being immediately included in the zero trees.) We shall denote by  $Z$  the set of arcs of  $T$  represented in the zero trees. In an abuse of notation, we shall occasionally use  $T$  to denote the set of arcs in the basis tree. Then  $T - Z$  is the set of basis arcs not represented in the zero trees.

We call a vertex  $i$  *eligible* if  $i$  is in the same zero tree as the root  $v$  of  $T$ . This definition is more restrictive than Orlin's ([19,20], Section 3) but is equivalent if all the zero arcs are in zero trees.

The algorithm maintains a set  $N^*$  that is a subset of the vertices whose potentials have not yet changed during the current phase. At the beginning of a phase,  $N^*$  is initialized to contain all the vertices. As the phase proceeds,  $N^*$  shrinks. The phase ends when  $N^*$  is empty.

A vertex  $i$  is *awake* if  $i \in N^*$  or if  $\pi_i$  is an integral multiple of  $\varepsilon/4$  ([19,20], Section 4). A nontree arc  $(i, j)$  is *admissible* for the  $\varepsilon$ -phase if vertex  $i$  is eligible and awake, and  $c_{ij}^\pi \leq -\varepsilon/4$ . Note that reversals of tree arcs are not admissible since they have nonnegative reduced costs. Again, our notion of admissibility is more restrictive than Orlin's, but this is only a technical detail to deal with the delayed inclusion of zero arcs in zero trees.

We associate with each vertex  $i$  three values in addition to its potential  $\pi_i$ . The first is a bit called the *flag* of  $i$  that is zero if  $\pi_i$  has not changed during the current phase and not all edges  $(i, j)$  have been examined for admissibility, and one otherwise. The flags represent  $N^*$ :  $i \in N^*$  if and only if  $\text{flag}(i) = 0$ . The second value for  $i$ , called its *mu-value*, is  $\min\{-c_{ki}^\pi \mid (k, i) \in T - Z\}$ . Note that  $c_{ki}^\pi \leq 0$  for  $(k, i) \in T$ , so all mu-values are nonnegative.

The third value for  $i$ , called its *offset*, is in general  $\varepsilon/4 - \pi_i \pmod{\varepsilon/4}$ , except in one special case: if  $\varepsilon/4 - \pi_i \pmod{\varepsilon/4} = \varepsilon/4$  and for the current value of  $\pi_i$  the algorithm has not yet finished examining  $i$  for outgoing admissible arcs, then the offset of  $i$  is zero.

We are now able to state a version of Orlin's scaling phase (procedure *improve-approximation* in [19,20], Section 4). At the beginning of the phase, all vertices have a flag of zero. The phase ends when all vertices have a flag of one. ( $N^* = \emptyset$ .) The phase consists of repeating steps until  $N^* = \emptyset$ . A single step within a phase proceeds as follows:

- (1) Let  $Y$  be the zero tree containing the root vertex  $v$ . Find a vertex  $i$  in  $Y$  with zero flag; if there is no such vertex, go to 2. Find an admissible arc  $(i, j)$  and pivot on it, ending the step; if there is no such arc, set the flag of  $i$  equal to one (deleting  $i$  from  $N^*$ ), and set the offset of  $i$  equal to  $\varepsilon/4 - \pi_i \pmod{\varepsilon/4}$ , ending the step.
- (2) (No vertex in  $Y$  has a zero flag.) Find the minimum mu-value  $\Delta_1$  and the minimum offset  $\Delta_2$  of the vertices in  $Y$ . If  $\Delta_1 < \Delta_2$ , go to 3; otherwise, go to 4.
- (3) ( $\Delta_1 < \Delta_2$ ). Add  $\Delta_1$  to the potentials of all vertices in  $Y$  and subtract  $\Delta_1$  from the offsets of all vertices in  $Y$ . Find an arc  $(k, i) \in T - Y$  with  $i \in Y$  and  $c_{ki}^\pi = 0$ . (Such an arc must exist after updating the potentials.) Link the zero trees containing  $k$  and  $i$  by adding arc  $(k, i)$ , ending the step.

- (4) ( $\Delta_2 \leq \Delta_1$ ). Add  $\Delta_2$  to the potential of all vertices in  $Y$  and subtract  $\Delta_2$  from the offsets of all vertices in  $Y$ . Find a vertex  $i$  in  $Y$  with offset zero. (Such a vertex must exist after updating the offsets.) Find an admissible arc  $(i, j)$  and pivot on it, ending the step. If there is no such arc, reset the offset of  $i$  to  $\varepsilon/4$ , ending the step.

We make a few comments on the differences between our implementation and Orlin's. We split the search for admissible arcs into two cases (substeps 1 and 4) to deal with the two conditions defining the "awake" state for a vertex  $i$ : substep 1 handles the case of  $i \in N^*$  and substep 4 handles the case of  $\pi_i$  an integral multiple of  $\varepsilon/4$ . In both substeps, vertex  $i$  is guaranteed to be eligible and awake; thus to find an arc on which to pivot it suffices to find an arc  $(i, j)$  with  $c_{ij}^\pi < -\varepsilon/4$ . The search for such arcs in 1 and 4 uses the *current arc* mechanism of Goldberg and Tarjan [11], as discussed by Orlin ([19,20], Section 4). In Orlin's implementation, all zero arcs are automatically in the zero trees. Our algorithm adds arcs one-at-a-time in substep 3, growing the tree  $Y$  until an admissible arc can be found or until a non-zero potential change creates a new zero arc.

Now we fill in some of the details of the implementation. The termination test for  $N^* = \emptyset$  can be performed efficiently (in constant time) by maintaining a count of vertices in  $N^*$ . A *pivot* on an arc  $(i, j)$  consists of adding the arc  $(i, j)$  to the basis tree  $T$ , deleting some arc  $(k, l)$  on the resulting cycle, changing the root of  $T$  to  $k$  or  $l$ , and reversing the direction of appropriate arcs in  $T$  to effect the rerooting. All the required operations, including determining the leaving arc  $(k, l)$ , are performed using the original dynamic tree data structure of Sleator and Tarjan [21,22], as spelled out in [23]. the time per pivot for these operations is  $O(\log n)$ .

We use the Euler tour structure to store the four values associated with each vertex, namely its potential, its flag, its mu-value, and its offset. (Thus we need four sets of the operations *find-val*, *find-min-val*, *change-val*, and *add-val*, although only a subset of the sixteen possible operations are actually used.) In order to maintain the mu-values, we also need to maintain for each vertex  $i$  a heap that contains the values  $-c_{ki}^\pi$  such that  $(k, i) \in T - Z$ . The minimum element in the heap for  $i$  is the mu-value for  $i$ . Values in each heap are stored in difference form. This allows adding an increment to all elements of a heap in constant time. We can use any standard heap implementation that supports the operations of insertion, deletion, and finding the minimum in time logarithmic in the heap size; see e.g. [25]. The article [9] includes a discussion of storing heap values in difference form.

Each execution of 1–4 takes a constant number of operations on the Euler tour structure and on the auxiliary heaps. (We leave the determination of exactly what operations are necessary as a routine exercise for the reader.) Thus the time required for one such execution is  $O(\log n)$ . Combining this with Orlin's bound of  $O((nm) \min\{\log(nC), m \log n\})$  on the total number of executions of 1–4 gives us a bound of  $O((nm \log n) \min\{\log(nC), m \log n\})$  on the total running time. Here  $n$ ,  $m$ , and  $C$  are the number of vertices, number of edges, and maximum absolute value of

an arc cost, respectively; the bound in terms of  $C$  requires integer arc costs. Our bound is better than Orlin's original bound by a factor of  $O(n/\log n)$ . As noted in the introduction, the same bound as ours can be obtained by using the original dynamic tree data structure extended as described in [10], but this structure is significantly more complicated than the Euler tour structure.

## References

- [1] G.M. Adel'son-Vel'skii and E.M. Landis, An algorithm for the organization of information, *Soviet Math. Dokl.* 3 (1962) 1259–1262.
- [2] R. Bayer and E. McCreight, Organization of large ordered indexes, *Acta Inform.* 1 (1972) 173–189.
- [3] S. Bent, D. Sleator and R.E. Tarjan, Biased search trees, *SIAM J. Computing* 14 (1985) 545–568.
- [4] R.F. Cohen and R. Tamassia, Dynamic expression trees and their applications, *Proc. 2nd ACM–SIAM Symposium on Discrete Algorithms* (1991) 52–61.
- [5] J. Edmonds and R.M. Karp, Theoretical improvements in algorithmic efficiency for network flow problems, *J. Assoc. Comput. Mach.* 19 (1972) 248–264.
- [6] D. Eppstein, G.F. Italiano, R. Tamassia, R.E. Tarjan, J. Westbrook and M. Yung, Maintenance of a minimum spanning forest in a dynamic planar graph, *Proc. 1st ACM–SIAM Symp. on Discrete Algorithms* (1990) 1–11.
- [7] G.N. Frederickson, Data structures for on-line updating of minimum spanning trees, *SIAM J. Comput.* 14 (1985) 781–798.
- [8] G.N. Frederickson, Ambivalent data structures for dynamic 2-edge-connectivity and  $k$  smallest spanning trees, *Proc. 32nd IEEE Symp. on Foundations of Computer Science* (1991) 632–641.
- [9] H.N. Gabow, Z. Galil, T. Spencer and R.E. Tarjan, Efficient algorithms for finding minimum spanning trees in undirected and directed graphs, *Combinatorica* 6 (1986) 109–122.
- [10] A.V. Goldberg, M.D. Grigoriadis, and R.E. Tarjan, Use of dynamic trees in a network simplex algorithm for the maximum flow problem, *Math. Prog.* 50 (1991) 277–290.
- [11] A.V. Goldberg and R.E. Tarjan, A new approach to the maximum flow problem, *J. Assoc. Comput. Mach.* 35 (1988) 921–940.
- [12] A.V. Goldberg and R.E. Tarjan, Finding minimum-cost circulations by canceling negative cycles, *J. Assoc. Comput. Mach.* 36 (1989) 873–886.
- [13] A.V. Goldberg and R.E. Tarjan, Finding minimum-cost circulations by successive approximation, *Math. of Oper. Res.* 15 (1990) 430–466.
- [14] D. Goldfarb and J. Hao, A primal simplex algorithm that solves the maximum flow problem in at most  $nm$  pivots and  $O(n^2m)$  time, *Mathematical Programming* 47 (1990) 353–365.
- [15] L.J. Guibas and R. Sedgwick, A dichromatic framework for balanced trees, *Proc. 19th Annual IEEE Symposium on Foundations of Computer Science* (1978) 8–21.
- [16] M.R. Henzinger and V. King, Randomized dynamic graph algorithms with polylogarithmic time per operation, *Proc. 27th Annual ACM Symp. on Theory of Computing* (1995) 519–527.
- [17] G.L. Miller and J.H. Reif, Parallel tree contraction and its application, *Proc. 26th Annual IEEE Symp. on Foundations of Comp. Sci.* (1985) 478–489.
- [18] P.B. Miltersen, S. Subramanian, J.S. Vitter and R. Tamassia, Complexity models for incremental computation, *Theoretical Computer Science* 130 (1994) 203–236.
- [19] J.B. Orlin, A polynomial time primal network simplex algorithm for minimum cost flows (an extended abstract), *Proc. 7th ACM–SIAM Symp. on Discrete Algorithms* (1996) 474–481.
- [20] J.B. Orlin, A polynomial time primal network simplex algorithm for minimum cost flows, *Mathematical Programming* 78 (1997) 109–129.
- [21] D. Sleator and R.E. Tarjan, A data structure for dynamic trees, *J. Computer and System Sciences* 26 (1983) 362–391.
- [22] D. Sleator and R.E. Tarjan, Self-adjusting binary search trees, *J. Assoc. Comput. Mach.* 32 (1985) 652–686.
- [23] R.E. Tarjan, Efficiency of the primal network simplex algorithm for the minimum-cost circulation problem, *Math. of Oper. Res.* 16 (1991) 272–291.



- [24] R.E. Tarjan, Updating a balanced search tree in  $O(1)$  rotations, *Information Processing Letters* 16 (1983) 253–257.
- [25] R.E. Tarjan, Data Structures and Network Algorithms, *CBMS 44*, Society for Industrial and Applied Mathematics, Philadelphia, PA, 1983.
- [26] R.E. Tarjan and U. Vishkin, An efficient parallel biconnectivity algorithm, *SIAM J. Comput.* 14 (1985) 862–874.