

# Acyclic Joins

Jef Wijsen

April 26, 2015

## 1 Motivation

### Joins in a Distributed Environment

Assume the following relations.<sup>1</sup>

- $M[NN, Field\_of\_Study, Year]$  stores data about students of UMONS. For example, (19950423158, Informatics, BAC3) states that the person with national number 19950423158 is enrolled in BAC3 Informatics. The relation  $M$  is stored in Mons.
- $B[NN, Street, Number, City]$  stores the addresses of all Belgian citizens. The relation  $B$  is stored in Brussels.

UMONS wants to get the join  $M \bowtie B$ . Several computations are possible.

1. Transmit relation  $B$  from Brussels to Mons, and compute the join at Mons. If we assume ten million Belgians and five thousand students, 99.95% of the transmitted tuples are *dangling*, meaning that they do not join with a tuple from  $M$ .
2. Transmit  $\pi_{NN}(M)$  (five thousand tuples) from Mons to Brussels. Compute the join  $B \bowtie \pi_{NN}(M)$  in Brussels, and transmit the result (five thousand tuples) to Mons. Finally, compute  $M \bowtie (B \bowtie \pi_{NN}(M))$  in Mons. In this way, only ten thousand tuples are transmitted.

### Order of Joins

Assume relations  $R[AB]$ ,  $S[BC]$ ,  $T[CD]$ , which now reside on a single site. Assume that there are no dangling tuples. We want to join the three relations. Several computations are possible.

1. First compute  $R \bowtie T$ , which contains  $|R| \times |T|$  tuples. Then compute  $S \bowtie (R \bowtie T)$ , which may contain much less than  $|R| \times |T|$  tuples.
2. First compute  $R \bowtie S$ , then  $T \bowtie (R \bowtie S)$ . Since there are no dangling tuples, it can be easily seen that the intermediate result will not be larger than the output relation.

### Questions

The following questions arise.

1. In a distributed join, can we minimize the amount of tuples transmitted?
2. If we have a join of more than two relations, can we join the relations in a way so as to minimize the size of the intermediate results?

---

<sup>1</sup>See the course *Bases de Données I* for definitions of relation and the operator  $\bowtie$ .

## 2 Preliminaries

We assume *relation names*  $R, S, R_1, S_1, R_2, S_2, \dots$ . Each relation name  $R$  is associated with a finite set of *attributes*, denoted  $\text{sort}(R)$ . Letters  $A, B, C, \dots$  denote attributes. We will write  $R[X]$  to denote that  $R$  is a relation name with  $\text{sort}(R) = X$ .

A *schema*  $\mathbf{S}$  is a finite set of relation names such that for all  $R_1, R_2 \in \mathbf{S}$ , if  $R_1 \neq R_2$ , then  $\text{sort}(R_1) \neq \text{sort}(R_2)$ . Thus, we require that no two distinct relation names are associated with the same set of attributes. This restriction is not fundamental, but simplifies the technical treatment: an element  $R[X]$  of a schema is uniquely identified by  $X$ .

A *database* over a schema  $\mathbf{S}$  associates to each relation name  $R \in \mathbf{S}$  a relation over  $\text{sort}(R)$ .

Whenever a database is fixed, we do not distinguish between the relation name  $R$  and the relation associated with  $R$ . For example, when we talk about the “join of  $R$  and  $S$ ,” we mean the join of the relations associated with  $R$  and  $S$ .

Also, we will use  $R$  as a shorthand for  $\text{sort}(R)$ . For example, we will write  $\pi_R(R \bowtie S)$  instead of  $\pi_{\text{sort}(R)}(R \bowtie S)$ .

## 3 Semijoin

Recall that  $\text{sort}(R \bowtie S) := \text{sort}(R) \cup \text{sort}(S)$  and  $R \bowtie S := \{t \mid t[R] \in R \text{ and } t[S] \in S\}$ . The operator  $\bowtie$  is commutative and associative.

The *semijoin* of  $R$  and  $S$ , denoted  $R \ltimes S$ , is the subset of  $R$  containing each tuple of  $R$  that joins with some tuple of  $S$ . Formally,  $R \ltimes S := \pi_R(R \bowtie S)$ . A tuple of  $R$  that does not belong to  $R \ltimes S$  is called *dangling*.

**Exercise 1** Show that  $R \ltimes S = R \bowtie \pi_{R \cap S}(S)$ .

Assume that  $R$  and  $S$  reside on different sites, and that we want to compute  $R \ltimes S$ . The amount of transmitted data must be minimized. We can ship  $S$  to the site of  $R$ . However, the expression of Exercise 1 tells us that it is sufficient to ship  $\pi_{R \cap S}(S)$  to the site of  $R$ .

## 4 Joining Two Relations Residing at Different Sites

Show the following.

$$\begin{aligned} R \bowtie S &= (R \ltimes S) \bowtie S & (1) \\ &= (S \ltimes R) \bowtie R & (2) \end{aligned}$$

Assume that  $R$  and  $S$  reside on different sites, and that we want to compute  $R \bowtie S$ . Equation (1) tells us that we can compute  $R \ltimes S$  as in Section 3, and ship the result to the site of  $S$ . In this way, we avoid the transmission of dangling tuples of  $R$ . In summary [1, p. 701],

1. Compute  $\pi_{R \cap S}(S)$  at the site of  $S$ .
2. Ship  $\pi_{R \cap S}(S)$  to the site of  $R$ .
3. Compute  $R \ltimes S$  at the site of  $R$ , using the fact that  $R \ltimes S = R \bowtie \pi_{R \cap S}(S)$ .
4. Ship  $R \ltimes S$  to the site of  $S$ .
5. Compute  $R \bowtie S$  at the site of  $S$ , using the fact that  $R \bowtie S = (R \ltimes S) \bowtie S$ .

There is a symmetric strategy, with  $R$  and  $S$  interchanged.

| $R$ | $A$ | $B$ | $S$ | $B$ | $C$ | $T$ | $C$ | $D$ |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
|     | 1   | 2   |     | 1   | 2   |     | 1   | 2   |
|     | 2   | 4   |     | 2   | 4   |     | 2   | 4   |
|     | 3   | 6   |     | 3   | 6   |     | 3   | 6   |
|     | 4   | 8   |     | 4   | 8   |     | 4   | 8   |

Figure 1: Three relations to be joined.

| $R$ | $A$ | $B$ | $S$ | $B$ | $C$ | $U$ | $C$ | $A$ |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
|     | $a$ | $b$ |     | $b$ | $c$ |     | $c$ | $d$ |
|     | $d$ | $e$ |     | $e$ | $f$ |     | $f$ | $a$ |

Figure 2: Three relations to be joined.

## 5 Joining Three or More Relations

Let  $\mathbf{db}$  be a database over schema  $\mathbf{S} = \{R_1, \dots, R_n\}$ . We say that a tuple  $t$  of  $R_i$  is *dangling* with respect to  $\mathbf{S}$  if  $t \notin \pi_{R_i}(R_1 \bowtie R_2 \bowtie \dots \bowtie R_k)$ . We can try to eliminate dangling tuples by applying semijoins as in Section 4. A *semijoin program* for  $\mathbf{S}$  is a sequence of commands

$$\begin{aligned}
R_{i_1} &:= R_{i_1} \ltimes R_{j_1}; \\
R_{i_2} &:= R_{i_2} \ltimes R_{j_2}; \\
&\vdots \\
R_{i_p} &:= R_{i_p} \ltimes R_{j_p};
\end{aligned}$$

This is called a *full reducer* for  $\mathbf{S}$  if for each database  $\mathbf{db}$  over  $\mathbf{S}$ , applying this program yields a database without dangling tuples.

**Example 1** Consider the database of Fig. 1 and the semijoin program

$$\begin{aligned}
R &:= R \ltimes S \\
S &:= S \ltimes T \\
T &:= T \ltimes S
\end{aligned}$$

The first step eliminates tuples (3, 6) and (4, 8) from  $R$ , and the second step does the same to  $S$ . The third step eliminates (1, 2) and (3, 6) from  $T$ . If we then take the join of the three relations, we find that the only tuple in the join  $R \bowtie S \bowtie T$  is (1, 2, 4, 8). That is, tuple (2, 4) is still dangling in  $R$  and  $T$ , and tuple (1, 2) is dangling in  $S$ . Thus, this semijoin program is not a full reducer.

**Example 2** A full reducer for the relations of Fig. 1 is

$$\begin{aligned}
S &:= S \ltimes R \\
T &:= T \ltimes S \\
S &:= S \ltimes T \\
R &:= R \ltimes S
\end{aligned}$$

We will show in Theorem 2 that this program eliminates dangling tuples from  $R$ ,  $S$ , and  $T$  independent of the initial values of these relations.




Figure 3: Cyclic hypergraph.




Figure 4: Acyclic hypergraph.  $ABC$  is an ear that can be removed in favor of  $ACE$ , because  $ABC \setminus ACE = B$  and  $B$  is unique to  $ABC$ .

**Example 3** Consider the database of Fig. 2. Notice the following.

$$\begin{aligned}
 R &= R \bowtie S \\
 R &= R \bowtie U \\
 S &= S \bowtie R \\
 S &= S \bowtie U \\
 U &= U \bowtie R \\
 U &= U \bowtie S
 \end{aligned}$$

Since  $R \bowtie S \bowtie U = \{\}$ , it is correct to conclude that there exists no full reducer for this schema.

Example 3 raises an important question: which schemas have a full reducer?

## 6 Acyclic Schemas

A *hypergraph* is a pair  $(V, E)$  where  $V$  is a set of vertexes and  $E$  is a family of distinct nonempty subsets of  $V$ , called *hyperedges*.

The hypergraph of schema  $\mathbf{S}$  is the pair  $(V, E)$  where  $V = \bigcup \{\text{sort}(R) \mid R \in \mathbf{S}\}$  and  $E = \{\text{sort}(R) \mid R \in \mathbf{S}\}$ .

Let  $E$  and  $F$  be two hyperedges, and suppose that the attributes of  $E \setminus F$  are *unique* to  $E$ ; that is, they appear in no hyperedge but  $E$ . Then we call  $E$  an *ear*, and we term the removal of  $E$  from the hypergraph in question *ear removal*. We sometimes say “ $E$  is removed in favor of  $F$ ” in this situation. As a special case, if a hyperedge intersects no other hyperedge, then that hyperedge is an ear, and we can remove that hyperedge by “ear removal.”

The *GYO-reduction* of a hypergraph is obtained by applying ear removal until no more removals are possible. A hypergraph is *acyclic* if its GYO-reduction is the empty hypergraph; otherwise it is *cyclic*.

A schema is acyclic if its hypergraph is acyclic; otherwise it is cyclic.

**Exercise 2** Show that the hypergraph of Fig. 3 is cyclic, and that the hypergraph of Fig. 4 is acyclic.

**Theorem 1** *The GYO-reduction of a hypergraph is unique, independent of the sequence of ear removals chosen.*

**Proof** Note that a potential removal is still possible if another removal is chosen. For example, suppose  $E_1$  could be removed in favor of  $E_2$ . That is, the vertices of  $E_1 \setminus E_2$  are unique to  $E_1$ . We distinguish two cases.

1. If we do an ear removal of a hyperedge other than  $E_2$ , we can still remove  $E_1$ .
2. Suppose we first remove  $E_2$  in favor of some  $E_3$ . It suffices to show  $E_1 \setminus E_3 \subseteq E_1 \setminus E_2$ , so  $E_1$  is still an ear and can be removed in favor of  $E_3$ . Suppose towards a contradiction that there exists a vertex  $N \in E_1 \setminus E_3$  such that  $N \notin E_1 \setminus E_2$ . Then  $N \in E_2 \setminus E_3$  and  $N \in E_1$  (thus,  $N$  is not unique to  $E_2 \setminus E_3$ ), contradicting the assumption that  $E_2$  was an ear that could be removed in favor of  $E_3$ .

This concludes the proof.  $\square$

**Theorem 2** *A schema is acyclic if and only if it has a full reducer.*

**Proof of the  $\implies$ -direction** The proof runs by induction on the cardinality of  $\mathbf{S}$ . Clearly, if  $|\mathbf{S}| = 1$ , then the empty semijoin program is a full reducer for  $\mathbf{S}$ . For the induction step, let  $\mathbf{S}$  be an acyclic schema with  $|\mathbf{S}| \geq 2$ . Let  $\mathcal{G}$  be the hypergraph of  $\mathbf{S}$ . Since  $\mathcal{G}$  is acyclic, we can assume an ear  $S_1$  that can be removed in favor of some hyperedge  $T_1$ . Let  $\mathcal{H}$  be the resulting hypergraph, which must be acyclic. By the induction hypothesis, we can assume a full reducer  $P_{\mathcal{H}}$  for  $\mathbf{S} \setminus \{S_1\}$ . Consider the following semijoin program (call it  $P_{\mathcal{G}}$ ).

$$\begin{array}{l} T_1 := T_1 \bowtie S_1; \\ \boxed{\text{all commands of } P_{\mathcal{H}}} \\ S_1 := S_1 \bowtie T_1; \end{array}$$

Let  $S_1, \dots, S_n$  be an ordering of  $\mathbf{S}$  corresponding to a sequence of ear removals in a GYO reduction. Since  $S_1$  could be removed in favor of  $T_1$ , it follows

$$\text{sort}(S_1) \cap \left( \bigcup_{i=2}^n \text{sort}(S_i) \right) \subseteq \text{sort}(T_1) \quad (3)$$

We need to show that  $P_{\mathcal{G}}$  is a full reducer for  $\mathbf{S}$ . That is, we need to show that no tuple is dangling with respect to  $\mathbf{S}$ . We distinguish between tuples from  $S_2, \dots, S_n$ , and tuples from  $S_1$ .

**For every  $i \in \{2, \dots, n\}$ , no tuple of  $S_i$  is dangling with respect to  $\mathbf{S}$ .** Let  $i \in \{2, \dots, n\}$  and let  $s_i \in S_i$ . The full reducer  $P_{\mathcal{H}}$  ensures that  $s_i$  is not dangling with respect to  $\{S_2, \dots, S_n\}$ . That is, there exists a tuple  $t \in S_2 \bowtie \dots \bowtie S_n$  such that  $t[S_i] = s_i$ . The command  $T_1 := T_1 \bowtie S_1$  of  $P_{\mathcal{G}}$  ensures that  $t[T_1]$  joins with some tuple  $s_1 \in S_1$ . Since  $T_1 \in \{S_2, \dots, S_n\}$  and by (3), the tuple  $s_1$  joins with  $t$ . It follows that  $s_i$  is not dangling with respect to  $\mathbf{S}$ . Note also that  $s_1$  is not removed by the last command of  $P_{\mathcal{G}}$ .

**No tuple of  $S_1$  is dangling with respect to  $\mathbf{S}$ .** Let  $s_1 \in S_1$ . The command  $S_1 := S_1 \bowtie T_1$  of  $P_{\mathcal{G}}$  ensures that  $s_1$  joins with some tuple  $t_1 \in T_1$ . The full reducer  $P_{\mathcal{H}}$  ensures that  $t_1$  is not dangling with respect to  $\{S_2, \dots, S_n\}$  (recall that  $T_1 \in \{S_2, \dots, S_n\}$ ). Thus, there exists  $t \in S_2 \bowtie \dots \bowtie S_n$  such that  $t[T_1] = t_1$ . By (3),  $s_1$  joins with  $t$ , hence  $s_1$  is not dangling with respect to  $\mathbf{S}$ .  $\square$

**Example 4** The following GYO-reduction shows that schema  $\mathbf{S} = \{R[AB], S[BC], T[CD]\}$  is acyclic.

1. Remove the ear  $R$  in favor of  $S$ .
2. In  $\{S[BC], T[CD]\}$ , remove the ear  $S$  in favor of  $T$ .
3. Remove the ear  $T$ .

A full reducer for  $\mathbf{S}$  is built “from the inside out.”

1. The empty semijoin program is a full reducer for  $\{T\}$ .
2. A full reducer for  $\{S, T\}$  is given by

$$\begin{aligned} T &:= T \bowtie S; \\ S &:= S \bowtie T; \end{aligned}$$

3. A full reducer for  $\{R, S, T\}$  is given by

$$\begin{aligned} S &:= S \bowtie R; \\ T &:= T \bowtie S; \\ S &:= S \bowtie T; \\ R &:= R \bowtie S; \end{aligned}$$

## 7 Order of Joins

Let  $\mathbf{S}$  be an acyclic database schema. Suppose we have applied a full reducer. We must now join all relations. Suppose we have removed  $S_1, S_2, \dots, S_n$  in that order. That is,  $S_1$  was the first ear removed,  $S_2$  was the second ear removed, and so on. Assume that for  $i \in \{1, \dots, n-1\}$ ,  $S_i$  was removed in favor of  $T_i \in \{S_{i+1}, \dots, S_n\}$ . In particular,  $T_{n-1} = S_n$ . The full reducer in the proof of Theorem 2 is the following.

$$\begin{aligned} T_1 &:= T_1 \bowtie S_1 \\ T_2 &:= T_2 \bowtie S_2 \\ &\vdots \\ T_{n-1} &:= T_{n-1} \bowtie S_{n-1} \\ S_{n-1} &:= S_{n-1} \bowtie T_{n-1} \\ &\vdots \\ S_i &:= S_i \bowtie T_i \\ &\vdots \\ S_2 &:= S_2 \bowtie T_2 \\ S_1 &:= S_1 \bowtie T_1 \end{aligned}$$

Now we join relations in reverse order, that is,

$$\begin{aligned} Result &:= S_n \\ Result &:= S_{n-1} \bowtie Result \\ Result &:= S_{n-2} \bowtie Result \\ &\vdots \\ Result &:= S_i \bowtie Result \\ &\vdots \\ Result &:= S_1 \bowtie Result \end{aligned}$$

We argue that the size of *Result* cannot decrease. When we join  $S_i$  to  $S_{i+1} \bowtie \dots \bowtie S_n$ , we know that every tuple of  $S_i$  joins with some tuple of  $S_{i+1} \bowtie \dots \bowtie S_n$ , because the command  $S_i := S_i \bowtie T_i$  in the full reducer

ensures that  $S_i$  has no dangling tuples with respect to  $\{S_{i+1}, \dots, S_n\}$ . As a consequence, no intermediate join can have more tuples than the output relation.

**Example 5** We continue Example 4. The GYO-reduction of  $\mathbf{S} = \{R[AB], S[BC], T[CD]\}$  shown there removes  $R, S, T$  in that order. So the order of the join is  $R \bowtie (S \bowtie T)$ . The command  $S := S \ltimes T$  in the full reducer (see Example 4) ensures that every tuple of  $S$  joins with some tuple of  $T$ . The command  $R := R \ltimes S$  in the full reducer ensures that every tuple of  $R$  joins with some tuple of  $S \bowtie T$ .

## References

- [1] J. D. Ullman. *Principles of Database and Knowledge-Base Systems, Volume II*. Computer Science Press, 1989.