Fast Nearest-neighbor Search in Disk-resident Graphs

Purnamrita Sarkar
Machine Learning Department
Carnegie Mellon University
psarkar@cs.cmu.edu

Andrew W. Moore
Google Inc.
Pittsburgh, PA
awm@google.com

ABSTRACT

Link prediction, personalized graph search, fraud detection, and many such graph mining problems revolve around the computation of the most "similar" k nodes to a given query node. One widely used class of similarity measures is based on random walks on graphs, e.g., personalized pagerank, hitting and commute times, and simrank. There are two fundamental problems associated with these measures. First, existing online algorithms typically examine the local neighborhood of the query node which can become significantly slower whenever high-degree nodes are encountered (a common phenomenon in real-world graphs). We prove that turning high degree nodes into sinks results in only a small approximation error, while greatly improving running times. The second problem is that of computing similarities at query time when the graph is too large to be memory-resident. The obvious solution is to split the graph into clusters of nodes and store each cluster on a disk page; ideally random walks will rarely cross cluster boundaries and cause page-faults. Our contributions here are twofold: (a) we present an efficient deterministic algorithm to find the k closest neighbors (in terms of personalized pagerank) of any query node in such a clustered graph, and (b) we develop a clustering algorithm (RWDISK) that uses only sequential sweeps over data files. Empirical results on several large publicly available graphs like DBLP, Citeseer and Live-Journal ($\sim$ 90 M edges) demonstrate that turning high degree nodes into sinks not only improves running time of RWDISK by a factor of 3 but also boosts link prediction accuracy by a factor of 4 on average. We also show that RWDISK returns more desirable (high conductance and small size) clusters than the popular clustering algorithm METIS, while requiring much less memory. Finally our deterministic algorithm for computing nearest neighbors incurs far fewer page-faults (factor of 5) than actually simulating random walks.

Categories and Subject Descriptors

H.3.3 [**Information Systems**]: Information Storage and Retrieval

General Terms

Algorithms, Experimentation

Keywords

graphs, random walks, link prediction, external memory

1. INTRODUCTION

A number of important real world applications (e.g. collaborative filtering in recommender networks, link prediction in social networks, fraud detection, and personalized graph search techniques) rely on finding nearest neighbors in large graphs, where "nearness" is defined using graph-theoretic measures of similarity. In keyword specific search in databases (entity relation graph of authors, papers and words) one wants to find other entities which are similar to the query words. In the context of spam detection in the webgraph, the goal is to find other webpages which are similar to spam pages. In recommender networks (bipartite graph of customers and products) we want to find the products (e.g. movies in the Netflix graph) that are most similar to the products a customer has already rated. The graph structure often conveys interesting intuition about how similar two nodes are. For example, two nodes are "close" if they have many common neighbors; more generally, two nodes are similar if there are many short paths between them.

The ensemble of paths between a pair of nodes is an essential ingredient for many widely used random walk based proximity measures, such as personalized pagerank (PPV) [14], hitting time and commute times [1], simrank [15], harmonic functions [26] etc. PPV has been used for content based search in large publication databases [3, 6]. The effectiveness of OBJECTRANK [3] (which combines PPV from words to obtain ranking for queries) was demonstrated through a relevance feedback survey. PPV has also been used for spam detection in the WWW [12, 16]. Hitting and commute times have been empirically shown to be useful for a variety of applications such as ranking in recommender networks [5], query suggestion [19], image segmentation [20] and robust multibody motion tracking [21].

In spite of the wide applicability, there are limitations to what we can do when graphs become enormous. Some algorithms, such as streaming algorithms [24], must make passes over the entire dataset to answer any query; this can be prohibitively expensive in online settings. Others perform clever preprocessing so that queries can be answered efficiently [10, 23]. However these algorithms store information which can be used for computing a *specific* similarity measure (e.g., personalized pagerank in [10]). This paper introduces analysis and algorithms which try to address the

scalability problem in a generalizable way: not specific to one kind of graph partitioning nor one specific proximity measure.

Another broad class of algorithms estimate the similarity between the query node and other nodes by examining local neighborhoods around the query node [4, 2, 6, 23]. The intuition is that in order to compute nearest neighbors, hopefully one would not need to look too far away from the query node. However, one fundamental computational issue with these techniques is the presence of very high degree nodes in the network. These techniques rely on updating one node's value by combining that of its neighbors; whenever a high degree node is encountered these algorithms have to examine a much larger neighborhood leading to severely degraded performance. Unfortunately, real-world graphs contain such high-degree nodes which, though few in number, are easily reachable from other nodes and hence are often encountered in random walks. Our first contribution is a simple transform of the graph (turning high degree nodes into sinks) that can mitigate the damage while having a provably bounded impact on accuracy. Indeed, we show that they *improve* accuracy in certain tasks like link prediction. We present two new results in this direction: a) a closed-form relationship between personalized pagerank and discounted hitting times, and b) the effect of making some high degree nodes into sinks on personalized page-rank. Note that the two together can be used to compute the effect of high degree nodes on discounted hitting times as well.

Another problem linked to large graphs is that algorithms can no longer assume that the entire graph can be stored in memory. In some cases, clever graph compression techniques can be applied to fit the graphs into main memory, but there are at least three settings where this might not work. First, social networks are far less compressible than Web graphs [8]. Second, decompression might lead to an unacceptable increase in query response time. Third, even if a graph could be compressed down to a gigabyte (comfortable main memory size in 2009) it is undesirable to keep it in memory on a machine which is running other applications, and in which there are occasional user queries to the graph. A good example of this third case is the problem of searching personal information networks [7], which integrates the user's personal information with information from the Web and hence needs to be performed on the user's own machine for privacy preservation [9].

Is there an intuitive representation of a disk-resident graph such that any random walk based measure can be easily computed from this representation? The obvious solution is to split the graph into clusters of nodes and store each cluster on a disk page; ideally random walks will rarely cross cluster boundaries and cause page-faults. This *clustered representation* can be used to quickly simulate random walks from any graph node, and by extension, any similarity measure based on random walks. Still, while simulations are computationally cheap, they have a lot of variation, and for some real-world graphs lacking well-defined clusters, they often lead to many page-faults. We propose a *deterministic* local algorithm guaranteed to return nearest neighbors in personalized pagerank from the disk-resident clustered graph. The same idea can also be used for computing nearest neighbors in hitting times. This is our second contribution.

Finally, we develop a fully external-memory clustering algorithm (RWDISK) that uses only sequential sweeps over data files. This serves as a preprocessing step that yields

the disk-resident clustered representation mentioned above, on top of which any nearest-neighbor algorithms can be run.

Extensive empirical results on several large publicly available graphs like DBLP, Citeseer and Live-Journal ($\sim$ 90 M edges) demonstrate that turning high degree nodes into sinks not only improves running time of RWDISK by a factor of 3 but also boosts link prediction accuracy by a factor of 4 on average. Our deterministic algorithm for computing nearest neighbors incurs far fewer page-faults (factor of 5) than actually simulating random walks. Finally, we show that RWDISK returns more desirable (high conductance and small size) clusters than the popular clustering algorithm METIS [17], while requiring much less memory.

The paper is organized as follows: in section 3.1 we theoretically show the effect of high degree nodes on personalized pagerank (PPV). We also show the same for discounted hitting times by presenting a new result which expresses discounted hitting times in terms of PPV. In section 3.2 we present a deterministic local algorithm to compute top k nodes in personalized pagerank using these clusters. In section 3.3 we describe our RWDISK algorithm for clustering a disk resident graph by only using sequential sweeps of files. We conclude with experimental results on large disk-resident Live Journal, DBLP and Citeseer graphs.

2. BACKGROUND AND RELATED WORK

In this section we will briefly describe interesting random walk based proximity measures, namely personalized pagerank and hitting times. We will also discuss the relevance of personalized pagerank for graph clustering.

Personalized Pagerank. Consider a random walk starting at node a, such that at any step the walk can be reset to the start node with probability α. The stationary distribution corresponding to this stochastic process is defined as the personalized pagerank vector (PPV) of node a. The entry corresponding to node j in the PPV vector for node a is denoted by $PPV(a, j)$. Large values of $PPV(a, j)$ is indicative of higher similarity/relevance of node j w.r.t a. For a general restart probability distribution $\boldsymbol{r}$ personalized pagerank is defined as $\boldsymbol{v} = \alpha \boldsymbol{r} + (1 - \alpha) P^T \boldsymbol{v}$. P is the row normalized probability transition matrix and $P^T \boldsymbol{v}$ is the distribution after one step of random walk from $\boldsymbol{v}$. Let's define $\boldsymbol{x_t}$ as the probability distribution over all nodes at timestep t. $\boldsymbol{x_0}$ is defined as the probability distribution with 1.0 at the start node and zero elsewhere. By definition we have $\boldsymbol{x_t} = P^T \boldsymbol{x_{t-1}} = (P^T)^t \boldsymbol{x_0}$. Let $\boldsymbol{v_t}$ be the partial sum of occupancy probabilities up to timestep t. Now we can write PPV as:

$$v(j) = \sum_{t=1}^{\infty} \alpha (1-\alpha)^{t-1} x_{t-1}(j) = \lim_{t \to \infty} v_t(j) \qquad (1)$$

Personalized pagerank has been shown to have empirical benefits in keyword search [6], link prediction [18], fighting spam [16]; there has been an extensive literature on algorithms for computing them locally [4, 6], off-line [14, 10, 23], and from streaming data [24] etc.

Discounted Hitting Time. Hitting time in random walks is a well-studied measure in probability theory [1]. Hitting times and other local variations of it has been used as a proximity measure for link prediction [18], recommender systems [5], query suggestion [19], manipulation resistant reputation systems [13] etc. We would use the following variation of hitting time. Note that this is closer to the original hitting time definition, and is different from the generalized hitting

time defined in [11]. Consider a random walk which, once started from i stops if node j is encountered, or with probability α. The expected time to hit node j in this process is defined as the α discounted hitting time from node i to node j, $(h_\alpha(i,j))$. Similar to the undiscounted hitting time, this can be written as the average of the hitting times of its neighbors to j. The definition is given by:

$$h_\alpha(i,j) = \begin{cases} 1 + (1-\alpha)\sum_k P(i,k)h_\alpha(k,j) & \text{when } i \neq j \\ 0 & \text{otherwise} \end{cases}$$
$$(2)$$

The maximum value of this quantity is $1/\alpha$, which happens when node j is never hit.

Graph Clustering. Recently there has been interesting theoretical work ([25, 2]) for using random walk based approaches for computing good quality local graph partitions (cluster) near a given anchor node. The main intuition is that a random walk started inside a low conductance cluster will mostly stay inside the cluster. Cluster-quality is measured by its conductance, which is defined as follows: For a subset of A of all nodes V, let $\Phi_V(A)$ denote conductance of A, and $\mu(A) = \sum_{i \in A} degree(i)$. As in [25], conductance is defined as:

$$\Phi_V(A) = \frac{E(A, V-A)}{min(\mu(A), \mu(V-A))} \qquad (3)$$

A good-quality cluster has small conductance, resulting from a small number of cross-edges compared to the total number of edges. The smaller the conductance the better the cluster quality. Hence 0 is perfect score, for a disconnected partition, whereas 1 is the worst score for having a cluster with no intra-cluster edges. Conductance of a graph is defined as the minimum conductance of all subsets A.

The formal algorithm to compute a low conductance local partition near a pre-selected seed node was given in [25]. The idea is to compute sparse representation of probability distribution over the neighboring nodes of a seed node in order to return a local cluster with small conductance with high probability. The running time is nearly linear in the size of the cluster it outputs.

3. PROPOSED WORK

There are two main problems with nearest neighbor computation in large real world networks. First, most local algorithms for computing nearest neighbors suffer from the presence of high degree nodes. In section 3.1 we propose a solution that converts high degree nodes to sinks. This effectively stops a random walk once it hits a high degree node, thus preventing a possible blow-up in the neighborhood-size in local algorithms. Our results imply that in power law graphs, this does not affect the proximity measures significantly.

The second issue is that of computing proximity measures on large disk-resident graphs. While there are existing external-memory algorithms for computing random walks in large disk-resident graphs [10, 23], most of these store sketches aimed to compute one particular measure. Streaming algorithms [24] on the other hand require multiple passes over the entire data. While all these algorithms use interesting theoretical properties of random walks, they do not provide a generic framework for computing *arbitrary* random walk based proximity measures on the fly. One solution would be to cluster the graph and store each partition on a disk-page. Given such a clustered representation,

one may easily simulate random walks, and thus compute nearest neighbors in hitting times, pagerank, simrank etc. Instead in section 3.2 we propose a deterministic local algorithm to compute nearest neighbors, which is later shown in section 4 to reduce number of page-faults compared to random simulations.

Finding a good clustering is a well-studied problem [2, 25]. Good clusters will have few cross edges, leading to self-contained random walks and less page-faults. Sometimes a good clustering can be achieved by using extra features, e.g. URL's in the web-graph. However we are unaware of a fully external memory clustering algorithm in the general setting. Building on some ideas in prior literature [2], [23] we present an external memory clustering algorithm in section 3.3.

3.1 Effect of High Degree Nodes

Local algorithms estimate the similarity between the query node and others by examining local neighbor-hoods around the query node. These mostly rely on dynamic programming or power iteration like techniques which involve updating a node's value by combining that of its neighbors. As a result whenever a high degree node is encountered these algorithms have to examine a much larger neighborhood leading to performance bottlenecks.

The authors in [23] use rounding in order to obtain sparse representations of personalized pagerank and simrank. However before rounding, the simrank/pagrank vectors can become very dense owing to the high degree nodes. In [4] and [6] the authors maintain a priority queue to store the active neighborhood of the query node. Every time a high degree node is visited all its neighbors need to be enqueued, thus slowing both algorithms down.

Because of the power-law degree distribution such high degree nodes often exist in real world networks. Although there are only a few of them, due to the small-world property these nodes are easily reachable from other nodes, and they are often encountered in random walks. We will discuss the effect of high degree nodes on two proximity measures, personalized pagerank and discounted hitting times. Our analysis of the effect of degree on hitting time, and personalized pagerank holds only for the case of undirected graphs, although we believe that the intuition is true for directed graphs as well. However, the main theorems, i.e. 3.5 and 3.1 hold for any graph.

Effect on Personalized Pagerank. The main intuition behind this analysis is that a very high degree node passes on a small fraction of its value to the out-neighbors, which might not be significant enough to invest our computing resources on. We argue that stopping a random walk at a high degree node does not change the personalized pagerank value at other nodes which have relatively smaller degree. First we show that the error incurred in personalized pagerank is inversely proportional to the degree of the sink node. Next we analyze the error for introducing a set of sink nodes. We turn a high degree node into a sink by removing all the out-going neighbors and adding one self-loop with probability one, to have a well-defined probability transition matrix P. We do not change any incoming edges.

We denote by $PPV(\boldsymbol{r}, j)$ the personalized pagerank value at node j w.r.t start distribution $\boldsymbol{r}$, and $PPV(i,j)$ denotes ppv value at node j w.r.t a random walk started at node i. Let $\widehat{PPV}$ be the personalized pagerank w.r.t. start distribution r on the changed transition matrix.

Theorem 3.1. *In a graph G, if a node s is changed into a sink, then for any node $i \neq s$, the personalized pagerank in the new graph w.r.t start distribution $\boldsymbol{r}$ can be written as:*

$$\widehat{PPV}(\boldsymbol{r}, i) = PPV(\boldsymbol{r}, i) - PPV(s, i)\frac{PPV(\boldsymbol{r}, s)}{PPV(s, s)}$$

Given theorem 3.1 we will prove that if degree of s is much higher that than of i, then the error will be small. In order to do this we would need to examine the quantity $PPV(i, j)/PPV(j, j)$. Define the first occurrence probability $fa_\alpha(i, j)$. Consider a random walk which stops if it hits node j; if j is not hit, it stops with probability α. $fa_\alpha(i, j)$ is simply the probability of hitting a node j for the first time from node i, in this α-discounted walk. This is defined as:

$$fa_\alpha(i, j) = \begin{cases} (1-\alpha)\sum_k P(i,k)fa_\alpha(j,k) & \text{when } i \neq j \\ 1 & \text{otherwise} \end{cases} \quad (4)$$

Lemma 3.2. *The personalized pagerank from node i to node j can be expressed as:*

$$PPV(i, j) = fa_\alpha(i, j) \times PPV(j, j)$$

PROOF SKETCH. As proven in [14, 10], the personalized pagerank from node i to node j is simply the probability that a length L path from i will end in j, where L is chosen from a geometric distribution with probability $P(L = t) = \alpha(1 - \alpha)^{t-1}$. The intuition is that, these paths can have multiple occurrences of j. If we condition on the first occurrence of j, then $PPV(i, j)$ would simply be probability of hitting j for the first time in an α discounted random walk times the personalized pagerank from j to itself.

The proof can be obtained by substituting eq. (4) in the recursive definition of PPV (as in eq. (5)). □

Lemma 3.3. *The error introduced at node $i \neq s$ by converting node s into a sink in an undirected graph, can be upper bounded by $\frac{d_i}{d_s}$.*

PROOF SKETCH. Note that by linearity of personalized pagerank vectors, we have $PPV(\boldsymbol{r}, i) = \sum_k \boldsymbol{r}_k PPV(k, i)$. Also as shown in the appendix, $PPV(s, i) \leq d_i PPV(i, s)/d_s \leq d_i/d_s$. Now, the above statement can be proved by combining linearity with $PPV(k, s)/PPV(s, s) = fa_\alpha(k, s) \leq 1$ (from lemma 3.2). □

Hence if d_s is much larger than d_i then this error is small. Now we will present the error for converting a set of nodes S to a sink. The first step is to show that the error incurred by turning a number of high degree nodes into sinks is upper bounded by the sum of their individual errors. That can again be simplified to the result in the following lemma.

Lemma 3.4. *In an undirected graph, if we convert all nodes in set $S = \{s_1, s_2, ..., s_k\}$ into sinks, then the error introduced at node $i \notin S$ can be upper bounded by $\dfrac{d_i}{\min_{s \in S} d_s}$.*

The proofs for theorem 3.1 and lemma 3.4 can be found in the appendix.

In real world networks the degree distribution often follows a power law, i.e. there are relatively fewer nodes with very large degree. And also most nodes have very low degree *relative* to these nodes. Hence we can make a few nodes

into sinks and gain a lot of computational efficiency without losing much accuracy.

Effect on Hitting Time. In order to see the effect of turning high degree nodes into sinks on discounted hitting times, we introduce the following result in this paper.

Theorem 3.5. *The α-discounted hitting time $h_\alpha(i, j)$ is related to personalized pagerank by:*

$$\boldsymbol{h}_\alpha(i, j) = \frac{1}{\alpha}\left[1 - \frac{PPV(i, j)}{PPV(j, j)}\right]$$

PROOF SKETCH. First we show that $h_\alpha(i, j) = \frac{1}{\alpha}(1 - fa_\alpha(i, j))$. This can be easily verified by substituting this in the equation for hitting time, i.e. eq. (2). This combined with lemma 3.2 gives the result. □

In this section we will only show the effect of deleting one high degree node on hitting time. The effect of removing a set of high degree nodes follows from the analysis of the last section and would be skipped for brevity. We denote by $\hat{h}_\alpha(i, j)$ the hitting time after we turn node s into a sink. By combining theorems 3.1 and 3.5, we can upper and lower bound the change in hitting time, i.e. $\triangle h_\alpha(i, j) = \hat{h}_\alpha(i, j) - h_\alpha(i, j)$. Using a little algebra we get,

Lemma 3.6. *If $d_j/d_s < \alpha$, then*

$$h_\alpha(i, j) - \frac{d_j}{\alpha(\alpha d_s - d_j)} \leq \hat{h}_\alpha(i, j) \leq h_\alpha(i, j) + \frac{d_j}{\alpha(\alpha d_s - d_j)}$$

We use the fact that $PPV(j, j) \geq \alpha$ and lemma 3.2 to obtain the above. Note that, if d_j/d_s is much smaller than α, then $\triangle h_\alpha(i, j) \approx \frac{d_j}{\alpha^2 d_s}$. This is generally the case for very large d_s, since α is usually set to be 0.1 or 0.2. This implies that turning a very high degree node into sink has a small effect on hitting time.

In this section we have given theoretical justification for changing the very-high degree nodes into sinks. We have shown its effects on two well known random walk based proximity measures. In the next two sections we would demonstrate algorithms to compute nearest neighbors on a clustered graph representation, and an external memory algorithm to compute clustering.

3.2 Nearest-neighbors on clustered graphs

Given a clustered representation, one can easily simulate random walks from a node, to obtain nearest neighbors in different proximity measures. While simulations are computationally cheap, they have a lot of variation, and for some real-world graphs they often lead to many page-faults, owing to the absence of well-defined clusters. In this section we discuss how to use the clusters for deterministic computation of nodes "close" to an arbitrary query. As the measure of "closeness" from i, we pick the degree-normalized personalized pagerank, i.e. $PPV(i, j)/d_j$. d_j is the weighted degree of node j. This is a truly personalized measure, in the sense that a popular node gets a high score only if it has a very high personalized pagerank value. We will use the degree-normalized personalized pagerank as a proximity measure for link prediction in our experiments as well.

We want to compute nearest neighbors in $PPV(i, j)/d_j$ from a node i. For an undirected graph, we have $PPV(i, j)/d_j = PPV(j, i)/d_i$ (appendix). Hence it is equivalent to computing nearest neighbors in personalized pagerank to a node i.

For an un-directed graph we can easily change these bounds to compute nearest neighbors in personalized pagerank *from* a node. For computing personalized pagerank *to a node*, we will make use of the dynamic programming technique introduced by [14] and further developed for computing sparse personalized pagerank vectors by [23]. For a given node i, the PPV from j to it, i.e. $PPV(j, i)$ can be written as

$$PPV^t(j, i) = \alpha\delta(i) + (1 - \alpha) \sum_{k \in nbs(j)} P(j, k)PPV^{t-1}(k, i)$$

$$(5)$$

Now let us assume that j and i are in the same cluster S. Hence the same equation becomes

$$PPV^t(j, i) = \alpha\delta(i) + (1 - \alpha) \left[\sum_{k \in nbs(j) \cap S} P(j, k)PPV^{t-1}(k, i) + \sum_{k \in nbs(j) \cap \bar{S}} P(j, k)PPV^{t-1}(k, i) \right]$$

Since we do not have access to $PPV^{t-1}(k), k \notin S$, we will replace it with upper and lower bounds. The lower bound is simply zero, i.e. we pretend that S is completely disconnected to the rest of the graph. A random walk from outside S has to cross the boundary of S, $\delta(S)$ to hit node i. Hence $PPV(k, i) = \sum_{m \in \delta(S)} Pr_\alpha(X_m | X_{\delta(S)})PPV(m, i)$, where X_m denotes the event that *Random walk hits node m before any other boundary node for the first time*, and the event $X_{\delta(S)}$ denotes the event that the *random walk hits the boundary* $\delta(S)$. Since this is a convex sum over personalized pagerank values from the boundary nodes, this is upper bounded by $\max_{m \in \delta(S)} PPV(m, i)$. Hence we have the upper and lower bounds as follows:

$$lb^t(j, i) = \alpha\delta(i) + (1 - \alpha) \sum_{k \in nbs(j) \cap S} lb^{t-1}(k, i)$$

$$ub^t(j, i) = \alpha\delta(i) + (1 - \alpha) \left[\sum_{k \in nbs(j) \cap S} P(j, k)ub^{t-1}(k, i) + \{1 - \sum_{k \in nbs(j) \cap S} P(j, k)\} \max_{m \in \delta(S)} ub^{t-1}(m, i) \right]$$

Since S is small in size, the power method suffices for computing these bounds, one could also use rounding methods introduced by [23]. At each iteration we maintain the upper and lower bounds for nodes within S, and at the global upper bound $\max_{m \in \delta(S)} ub^{t-1}(m, i)$. In order to expand S we bring in the clusters for x of the *external* neighbors of $\arg\max_{m \in \delta(S)} ub^{t-1}(m, i)$. Once this *global upper bound* falls below a pre-specified small threshold γ, we use these bounds to compute approximate k closest neighbors in degree-normalized personalized pagerank.

The ranking step to obtain top k nodes using upper and lower bounds is simple: we return all nodes which have lower bound greater than the $(k+1)^{th}$ largest upper bound (when $k = 1$, k^{th} largest is the largest probability). We denote this as ub_{k+1}. Since all nodes outside the cluster are guaranteed to have personalized pagerank smaller than the global upper bound, which in turn is smaller than γ, we know that the true $(k+1)^{th}$ largest probability will be smaller than ub_{k+1}. Hence any node with lower bound greater than ub_{k+1} is guaranteed to be greater than the $(k + 1)^{th}$ largest probability. We use an additive slack, e.g. $(ub_{k+1} - \epsilon)$ in order

to return the top k *approximately* large PPV nodes. The reason for using an additive slack is that, for larger values of ub_{k+1}, this behaves like a small relative error, whereas for small ub_{k+1} values it allows a large relative slack, which is useful since we do not want to spend energy on the tail of the rank list anyways. In our algorithm we initialize γ with 0.1 and keep decreasing it until the bounds are tight enough to return k largest nodes. Note that one could rank the probabilities using the lower bounds, and return top k of those after expanding the cluster a fixed number of times. This translates to a larger approximation slack.

Consider the case where we want to use this algorithm on a graph with high degree nodes converted into sinks. Since the altered graph is not undirected anymore, the ordering obtained from the degree normalized PPV from node i can be different from the ordering using PPV *to node* i. But using our error bounds from section 3.1 we can easily show that (skipped for brevity) if the difference between two personalized pagerank values $\widehat{PPV}(a, i)$ and $\widehat{PPV}(b, i)$ is larger than $\frac{2d_i}{\min_{s_i \in S} d(s_i)}$, then a will have larger degree-normalized PPV value than b, i.e. $\widehat{PPV}(i, a)/d_a \geq \widehat{PPV}(i, b)/d_b$.

Given that the networks follow a power law degree distribution, the minimum degree of the nodes made into sinks is considerably larger than d_i for most i, we see that the pairs which had a considerable gap in their PPV value *to i* should still have the same ordering in degree normalized PPV *from i*. Note that for high degree nodes the ordering will have more error. However because of the expander like growth of the neighborhood of a high degree node, most nodes are far away from it leading to an uninteresting set of nearest neighbors anyways.

3.3 Clustered Representation on Disk

So far we have discussed how to use a given clustered representation for computing nearest neighbors efficiently. Now we will present an algorithm to generate such a representation on disk. The intuition behind this representation is to use a set of anchor nodes and assign each remaining node to its "closest" anchor. Since personalized page-rank has been shown to yield good quality clusters [2], we use it as the measure of "closeness". Our algorithm starts with a random set of anchors, and compute personalized pagerank from them to the remaining nodes. Since all nodes might not be reached from this set of anchors, we iteratively add new anchors from the set of unreachable nodes, and the recompute the cluster assignments. Thus our clustering satisfies two properties: new anchors are far away from the existing anchors, and when the algorithm terminates, each node i is guaranteed to be assigned to its closest anchor. Even though the anchors are chosen randomly this should not affect the clustering significantly because, any node within a tight cluster can serve as the anchor for that cluster.

While clustering can be done based on personalized pagerank *from* or *to* a set of anchor nodes, one is not known to be better or worse than the other a priori. We use personalized pagerank *from the anchor nodes* as the measure of closeness. In [23] the authors presented a semi-external memory algorithm to compute personalized pagerank to a set of nodes. While the authors mention that their algorithm can be implemented purely via external memory manipulations, we close the loop via external memory computation of personalized pagerank from a set of anchor nodes (algorithm RWDISK). While we also use the idea of rounding introduced by the authors, the analysis of approximation error is

different from their analysis. We provide proof sketches for the main results in this section, details of the proof can be found in [22].

RWDISK. Personalized pagerank can be naively computed by power iterations. In eq. (1) there are two variables $\boldsymbol{x_t}$ and $\boldsymbol{v_t}$. While $x_t(i)$ is the probability of being at node i at time t, $v_t(i)$ is the cumulative sum of these occupancy probabilities with geometric weights. The RWDISK algorithm will sequentially read and write from four kinds of files. We will first introduce some notation. X_t denotes a random variable, which represents the node the random walk is visiting at time t. $x_0(a) = P(X_0 = a)$ is the probability that the random walk started at node a.

The *Edges* file remains constant. Each line represents an edge by a triplet $\{src, dst, p\}$, where `src` and `dst` are strings representing nodes in the graph, and $p = P(X_t = dst|X_{t-1} = src)$. *Edges* is sorted lexicographically by `src`.

At iteration t, the *Last* file contains x_{t-1}. Thus each line in *Last* is $\{src, anchor, value\}$, where `value` equals $P(X_{t-1} = src|X_0 = anchor)$.

At iteration t, the file *Newt* contains x_t, i.e. each line is $\{src, anchor, value\}$, where `value` equals $P(X_t = src|X_0 = anchor)$. Needless to say, the *Newt* of iteration t becomes the *Last* of iteration $t + 1$.

At iteration t of the algorithm, the file *Ans* represents the values for v_t. Thus each line in *Ans* is $\{src, anchor, value\}$, where $value = \sum_{n=1}^{t} \alpha(1-\alpha)^{n-1} x_{n-1}(j)$.

All of *Last*, *Newt*, and *Ans*, once construction is finished, will also be sorted by source.

Note that *Newt* is simply a matrix-vector product between the transition matrix stored in *Edges* and *Last*. Since both these files are sorted lexicographically, this can be obtained by a file-join like algorithm. The first step simply joins the two files, and accumulates the probability values at a node from its in-neighbors. In the next step the *Newt* file is sorted and compressed, in order to add up the values from different in-neighbors. We provide the details in [22]. Once we have *Newt*, we multiply the probabilities by $\alpha(1-\alpha)^{t-1}$ (the probability that a random walk will stop at timestep t, if at any step the probability of stopping is α) and accumulate the values into the previous *Ans* file. At the end of one iteration *Newt* file is renamed to *Last* file. We fix the number of iterations at `maxiter`. Here is the error bound for stopping at `maxiter`. This follows from the fact that the contribution from the later iterations decay with a factor of $1 - \alpha$.

Theorem 3.7. *Let $\boldsymbol{v_t}$ be the partial sum of distributions up to time t. If we stop at $t = $ `maxiter`, then the error is given by $\boldsymbol{v} - \boldsymbol{v_{maxiter}} = (1-\alpha)^{maxiter} PPV_\alpha(\boldsymbol{x_{maxiter}})$. where $PPV_\alpha(\boldsymbol{x_{maxiter}})$ is the personalized pagerank with start distribution $\boldsymbol{x_{maxiter}}$.*

This theorem indicates that the total error incurred by early stopping is $(1 - \alpha)^{maxiter}$.

One major problem is that intermediate files can become much larger than the number of edges (table 3.3). This is because, in most real-world networks within 4-5 steps it is possible to reach a huge fraction of the whole graph. Let N and E be the number of nodes and edges in the graph. So, the *Last* file can have at most $O(AN)$ lines (if all nodes are reachable from all anchors). If roughly n_a nodes can fit per page (our goal is to create a pagesize cluster for each anchor), then we would need N/n_a anchors. Hence the naive file join approach will lead to intermediate files of size $O(N^2/n_a)$. Since these files are also sorted in every iteration, this will

greatly affect both the runtime and disk-storage. This is why we introduce rounding.

Rounding for reducing file sizes. At timestep t only the values larger than ϵ_t are copied from *Last* to *Newt*. Elements of a sparse rounded probability density vector sums to at most one. Hence the total number of nonzero entries post-rounding can be at most $1/\epsilon_t$. As *Last* stores rounded x_{t-1} values for A anchors, its length can never exceed A/ϵ_{t-1}. Since *Newt* spreads the probability mass along the out-neighbors of each node in *Last*, its total length can be roughly $A \times d/\epsilon_{t-1}$, where d is the average out-degree. For $A = N/n_a$ anchors, this value is $E/(n_a \times \epsilon_{t-1})$. Without rounding this length was N^2/n_a.

We avoid sorting a huge internal *Newt* every iteration by updating the *Newt* in a *lazy* fashion. The details of this can be found in [22]. How big can the *Ans* file get? Since each *Newt* can be as big as $A \times d/\epsilon_{t-1}$, and we iterate for *maxiter* times, the size of *Ans* can be roughly as big as `maxiter` $\times A \times d/\epsilon_{t-1}$. Since the probability of stopping decreases by a factor of $(1-\alpha)^t$ with t, we gradually increase ϵ_t, leading to sparser solutions. Hence, $\epsilon_t \geq \epsilon, \forall t$. The file-sizes obtained from RWDISK with and without rounding is presented in table 3.3.

Algorithm	A	*Last* Size	*Newt* Size	*Ans* Size
No-rounding	$\frac{N}{n_a}$	$O(\frac{N^2}{n_a})$	$O(\frac{N^2}{n_a})$	$O(\frac{N^2}{n_a})$
Rounding	$\frac{N}{n_a}$	$O(\frac{N}{\epsilon \times n_a})$	$O(\frac{E}{\epsilon \times n_a})$	$O(\frac{maxiter \times E}{\epsilon \times n_a})$

Approximation Error. The rounding step saves vast amounts of time and intermediate disk space, but how much error does it introduce? In order to avoid confusion, we will use different notation (from section 3.1) for personalized pagerank before ($\boldsymbol{v}$) and after rounding ($\boldsymbol{\hat{v}}$).

Theorem 3.8. *Let $\boldsymbol{E}$ denote $\boldsymbol{v} - \boldsymbol{\hat{v}}$, and $PPV_\alpha(\boldsymbol{r})$ denote the personalized pagerank vector for start distribution $\boldsymbol{r}$, and restart probability α. We have $\boldsymbol{E} \leq \frac{1}{\alpha}PPV_\alpha(\boldsymbol{\bar{\epsilon}})$ where $\boldsymbol{\bar{\epsilon}}$ is a vector, with maximum entry smaller than $\frac{\alpha\epsilon}{1-\sqrt{1-\alpha}}$, and sum of entries less than 1.*

PROOF SKETCH. Denote $\widehat{\boldsymbol{x_t}}$ as the approximated $\boldsymbol{x_t}$ value from the rounding algorithm at iteration t. Let $\boldsymbol{E_t} = \boldsymbol{x_t} - \widehat{\boldsymbol{x_t}}$ be the error in the t^{th} step. $\widehat{\boldsymbol{x_t}}$ is strictly smaller than $\boldsymbol{x_t}$. Let $\boldsymbol{\bar{\epsilon}_t}$ be the vector of errors from rounding at time t. Note that, any entry of $\boldsymbol{\bar{\epsilon}_t}$ is at most ϵ_t. We can prove that $\boldsymbol{E_t} \leq \boldsymbol{\bar{\epsilon}_t} + P^T\boldsymbol{E_{t-1}} \leq \sum_{r=1}^{t}(P^T)^{t-r}\boldsymbol{\bar{\epsilon}_r}$.

Also, from the iterative definition of v (eq. 1) we have: $\boldsymbol{v} - \boldsymbol{\hat{v}} = \alpha \sum_{t=1}^{\infty}(1-\alpha)^{t-1}\boldsymbol{E_t}$. Combining the $\boldsymbol{E_t}$ vectors with the fact that $\epsilon_t \leq \epsilon_{t-1}/\sqrt{1-\alpha}$, gives us the result in theorem 3.8. $\square$

In the special case of undirected graphs theorem 3.8 implies that the error at any node is proportional to its degree.

Lemma 3.9. *The rounding error at any node i can be at most $\frac{d_i}{\delta} \frac{\epsilon}{1 - \sqrt{1-\alpha}}$, where d_i is the weighted degree of node i, and δ is the minimum weighted degree.*

For the proof of this, see the appendix. Combining the errors from theorems 3.8 and 3.7 we have the following expression of the total error.

Theorem 3.10 (Total Error). *If $\hat{v}_{maxiter}$ is obtained from RWDISK with parameters ϵ, α, `maxiter` and start distribution r then*

$$v - \hat{v}_{maxiter} \leq \frac{1}{\alpha} PPV_\alpha(\bar{\epsilon}) + (1 - \alpha)^{maxiter} PPV_\alpha(x_{maxiter})$$

where $x_{maxiter}$ is the probability distribution after `maxiter` steps, when the start distribution is r.

In [2], the authors adaptively pick the node with a significant amount of probability and propagate it to its neighbors. Our technique is similar in the sense that we spread the probabilities above a threshold in batch fashion. This is why, the form of the error bounds we get in theorem 3.8 is similar to [2], although the analysis is different. We would also like to point out that the authors used a lazy random walk. Although our random walks are not lazy, it can be shown using derivations from [2] that the personalized pagerank arising from an *un-lazy* random walk with restart probability α is equivalent to that with a lazy random walk with restart probability $\alpha/(2 - \alpha)$.

High Degree Nodes. In spite of rounding one problem with RWDISK is that if a node has high degree then it has large personalized pagerank value from many anchors and as a results can appear in a large number of $\{$`node,anchor`$\}$ pairs in the *Last* file. After the matrix multiplication each of these pairs now will lead to $\{$`nb,anchor`$\}$ pairs for each outgoing neighbor of the high degree node. Since we can only prune once the entire *Newt* file is computed the size can easily blow up. This is why RWDISK benefits from turning high degree nodes into sinks as described before in section 3.1.

4. RESULTS

We present our results in three steps: first we show the effect of high degree nodes on i) computational complexity of RWDISK, ii) page-faults in random walk simulations for an actual link prediction experiment on the clustered representation, and iii) link prediction accuracy. Second, we show the effect of deterministic algorithms for nearest-neighbor computation on reducing the total number of page-faults by fetching the right clusters. Last, we compare the usefulness of the clusters obtained from RWDISK w.r.t a popular in-memory algorithm METIS.

Data and System Details. We present our results on three of the largest publicly available social and citation networks: a connected subgraph of the Citeseer co-authorship network, the entire DBLP corpus, and Live Journal (table 1). We use an undirected graph representation, although RWDISK can be used for directed graphs as well. The experiments were done on an off-the-shelf PC. We used a size 100 buffer and the *least recently used* replacement scheme. Each time a random walk moves to a cluster not already present in the buffer, the system incurs page-faults. We used a pagesize of $4KB$, which is standard in most computing environments. This size can be much larger (e.g. $1MB$) in advanced architectures, and for those cases these experiments will have to be run with a re-tuned parameter setting.

4.1 Effect of High Degree Nodes

Turning high degree nodes into sinks have three-fold advantage: first, it drastically speeds up our external memory clustering; second, it reduces number of page-faults in random walk simulations done in order to rank nodes for link-

Dataset	Size of *Edges*	Nodes	Edges	Median Degree
Citeseer	$24MB$	$100K$	$700K$	4
DBLP	$283MB$	$1.4M$	$12M$	5
LiveJournal	$1.4GB$	$4.8M$	$86M$	5

Table 1: # nodes,directed edges in graphs

prediction experiments; second it actually *improves* link-prediction accuracy.

Dataset	Sink Nodes		Time
	Min degree	Number	
Citeseer	None	0	1.3 hours
DBLP	None	0	≥ 2.5 days
	1000	900	**11** hours
LiveJournal	1000	950	60 hours
	100	$134,000$	**17** hours

Table 2: For each dataset, the minimum degree, above which nodes were turned into sinks, and the total number of sink nodes, time for RWDISK.

Effect on RWDISK. Table 2 contains running times of RWDISK on three graphs. For Citeseer, RWDISK algorithm completed roughly in an hour without introducing any sink nodes. For DBLP, without degree-deletion, the experiments ran for above 2.5 days, after which they were stopped. After turning nodes with degree higher than 1000, the time was reduced to 11 hours, a larger than 5.5 fold speedup. The Live-Journal graph is the largest and most dense of all three. After we made nodes of degree higher than 1000 into sinks, the algorithm took 60 hours, which was reduced to 17 ($\geq$ 3 fold speedup) after removing nodes above degree 100. In table 1 note that, for both DBLP and LiveJournal, the median degree is much smaller than the minimum degree of nodes converted into sink nodes. This combined with our analysis in section 3.1 confirms that we did achieve a huge computational gain without sacrificing the quality of approximation.

Link Prediction. For link-prediction we use degree normalized personalized pagerank as the proximity measure for predicting missing links. We picked the same set of 1000 nodes and the same set of links from each graph before and after turning the high degree nodes into sinks. For each node i we held out $1/3^{rd}$ of its edges and reported the percentage of held-out neighbors in top 10 ranked nodes in degree-normalized personalized pagerank from i. Only nodes below degree 100 and above degree 3 were candidates for link deletion, so that no sink node can ever be a candidate. From each node 50 random walks of length 20 were executed. Note that this is not AUC score; so a random prediction does much worse that 0.5 in these tasks.

From table 3 we see that turning high-degree nodes into sinks not only decrease page-faults by a factor of ~ 7, it also boosts the link prediction accuracy by a factor of 4 on average. The fact that page-faults decrease after introducing sink nodes is obvious, since in the original graph every time a node hits a high degree node there is higher chance of incurring page-faults.

We believe that the link prediction accuracy is related to quality of clusters, and transitivity of relationships in a

Dataset	Sink nodes	Accuracy	Page-faults
LiveJournal	none	0.2	1502
	degree above 100	**0.43**	**255**
DBLP	none	0.1	1881
	degree above 1000	**0.58**	**231**
Citeseer	none	0.74	69
	degree above 100	0.74	67

Table 3: Mean link-pred. acc. and pagefaults

graph. More specifically in a *well-knit* cluster, two connected nodes do not just share one edge, they are also connected by many short paths, which makes link-prediction easy. On the other hand if a graph has a more expander-like structure, then in random-walk based proximity measures, everyone ends up being far away from everyone else. This leads to poor link prediction scores. In table 3 one can catch the trend of link prediction scores from worse to better from LiveJournal to Citeseer.

Our intuition about the relationship between cluster quality and predictability is reflected in figure 1, where we see that LiveJournal has worse page-fault/conductance scores than DBLP, which in turn has worse scores than Citeseer. Within each dataset, we see that turning high degree nodes into a sink generally helps link prediction, which is probably because it also improves the cluster-quality. Are all high-degree nodes harmful? In DBLP high degree nodes without exception end up being words which can confuse random walks. However the Citeseer graph only contains author-author connections, which are much less ambiguous than paper-word connections. There are also not as many high degree nodes as compared to the other datasets. This might be the reason why introducing sink nodes does not change the link prediction accuracy.

4.2 Deterministic vs. Simulations

We present the mean and median number of pagefaults incurred by the deterministic algorithm in section 3.2. We executed the algorithm for computing top 10 neighbors with approximation slack 0.005 for 500 randomly picked nodes. For Citeseer we computed the nearest neighbors in the orig-

Dataset	Mean Page-faults	Median Page-faults
LiveJournal	64	29
DBLP	54	16.5
Citeseer	6	2

Table 4: Page-faults for computing 10 nearest neighbors using lower and upper bounds

inal graph, whereas for DBLP we turned nodes with degree above 1000 into sinks and for LiveJournal we turned nodes with degree above 100 into sinks. Both mean and median pagefaults decrease from LiveJournal to Citeseer, showing the increasing cluster-quality, as is evident from the previous results. The difference between mean and median reveals that for some nodes the neighborhood is explored much more in order to compute the top 10 nodes. Upon closer investigation we found that for high degree nodes, the clusters have a lot of boundary nodes and hence the bounds are hard to tighten. Also from high degree nodes all other nodes are more or less farther away. In contrast to random simulations

(table 3), these results show the superiority of the deterministic algorithm over random simulations in terms of number of page-faults (roughly 5 fold improvement).

4.3 RWDISK vs. METIS

We used `maxiter` = 30, $\alpha = 0.1$ and $\epsilon = 0.001$ for PPV computation. We use PPV and RWDISK interchangeably in this section. Note that $\alpha = 0.1$ in our random-walk setting is equivalent to a restart probability of $\alpha/(2 - \alpha) = 0.05$ in the lazy random walk setting of [2].

We used METIS as a baseline algorithm [17][1], which is a state of the art *in memory* graph partitioning algorithm. We used METIS to break the DBLP graph into about $50,000$ parts, which used **20 GB** of RAM, and LiveJournal into about $75,000$ parts which used **50 GB** of RAM. Since METIS was creating comparably larger clusters we tried to divide the Live Journal graph into $100,000$ parts, however the memory requirement was **80 GB** which was prohibitively large for us. In comparison RWDISK can be executed on a $2 - 4$ GB standard computing unit. Table 1 contains the details of the three different graphs and table 2 contains running times of RWDISK on these. Although the clusters are computed after turning high degree nodes into sinks, the comparison with METIS is done on the original graphs.

Measure of cluster quality. A good disk-based clustering must combine two characteristics: (a) the clusters should have low conductance, and (b) they should fit in disk-sized pages. Now, the graph conductance ϕ measures the average number of times a random walk can escape outside a cluster [25], and each such escape requires the loading of one new cluster, causing an average of m page-faults ($m = 1$ if each cluster fits inside one page). Thus, $\phi \cdot m$ is the average number of page-faults incurred by one step of a random walk; we use this as our overall measure of cluster quality. Note that m here is the expected size (in pages) of the cluster that a randomly picked node belongs to, and this is *not* necessarily the average number of pages per cluster.

Briefly, figure 1 tells us that in a single step random walk METIS will lead to similar number of pagefaults on Citeseer, about $1/2$ pagefaults more than RWDISK on DBLP and 1 more in *LiveJournal*. Hence in a 20 step random walk METIS will lead to about 5 more pagefaults than RWDISK on DBLP and 20 more pagefaults on LiveJournal. Note that since a new cluster can be much larger than a disk-page size it is possible to make more than 20 pagefaults on a 20 step random walk in our paradigm. In order to demonstrate the accuracy of this measure we actually simulated 50 random walks of length 20 from 100 randomly picked nodes from the three different graphs. We noted the average page-faults and average time in wall-clock seconds. Figure 2 shows how many *more pagefaults METIS incurs than RWDISK in every simulation*. The wallclock seconds is the **total time** taken for all 50 simulations averaged over the 100 random nodes. These numbers exactly match our expectation from figure 1. We see that on Citeseer METIS and RWDISK give comparable cluster qualities, but on DBLP and LiveJournal RWDISK performs much better.

5. CONCLUSION

This paper address the following problem. Random-walk based measures of proximity in graphs, such as Personalized Page Rank, Hitting Times and Commute times are

---

[1]The software for partitioning power law graphs has not yet been released.

Figure 1: The histograms for the expected number of pagefaults if a random walk stepped outside a cluster for a randomly picked node. Left to right the panels are for Citeseer, DBLP and LiveJournal.


Figure 2: #Page-faults(METIS)-#Page-faults(RWDISK) per 20 step random walk in upper panel. Bottom Panel contains total time for simulating 50 such random walks. Both are averaged over 100 randomly picked source nodes.

becoming very important and popular, and yet there are limitations to what we can do when graphs become enormous. This paper introduces analysis and algorithms which try to address this in a generalizable way: not specific to one kind of graph partitioning nor one specific proximity measure. We take two steps. First, we identify the serious role played by high degree nodes in damaging computational complexity, and we prove that a simple transform of the graph can mitigate the damage with bounded impact on accuracy. Second, we apply the result to produce algorithms for the two components of general-purpose proximity queries on enormous graphs: algorithms to rank top-n neighbors by a broad class of random-walk based proximity measures including PPV, and a graph partitioning step to distribute graphs over a file system or over nodes of a distributed compute-node cluster. In future work we will experiment with a highly optimized implementation designed to respect true disk page size and hope to give results on graphs with billions of edges.

6. REFERENCES

[1] D. Aldous and J. A. Fill. *Reversible Markov Chains.* 2001.

[2] R. Andersen, F. Chung, and K. Lang. Local graph partitioning using pagerank vectors. In *FOCS*, 2006.

[3] A. Balmin, V. Hristidis, and Y. Papakonstantinou. ObjectRank: Authority-based keyword search in databases. In *VLDB, 2004.*

[4] P. Berkhin. Bookmark-Coloring Algorithm for Personalized PageRank Computing. *Internet Mathematics*, 2006.

[5] M. Brand. A Random Walks Perspective on Maximizing Satisfaction and Profit. In *SIAM '05*, 2005.

[6] S. Chakrabarti. Dynamic personalized pagerank in entity-relation graphs. In *WWW '07*, New York, NY, USA.

[7] S. Chakrabarti, J. Mirchandani, and A. Nandi. Spin: searching personal information networks. In *SIGIR '05*.

[8] F. Chierichetti, R. Kumar, S. Lattanzi, M. Mitzenmacher, A. Panconesi, and P. Raghavan. On compressing social networks. In *KDD '09*.

[9] B. B. Dalvi, M. Kshirsagar, and S. Sudarshan. Keyword search on external memory data graphs. *Proceedings of VLDB Endowment*, 1(1):1189–1204, 2008.

[10] D. Fogaras, B. Rcz, K. Csalogány, and T. Sarlós. Towards scaling fully personalized pagerank: Algorithms, lower bounds, and experiments. *Internet Mathematics*, 2004.

[11] F. C. Graham and W. Zhao. Pagerank and random walks on graphs. In *"Fete of Combinatorics"*.

[12] Z. Gyongyi, H. Garcia-Molina, and J. Pedersen. Combating web spam with trustrank. In *VLDB*, pages 576–587, 2004.

[13] J. Hopcroft and D. Sheldon. Manipulation-resistant reputations using hitting time. Technical report, Cornell University, 2007.

[14] G. Jeh and J. Widom. Scaling personalized web search. In *Stanford University Technical Report*, 2002.

[15] G. Jeh and J. Widom. Simrank: A measure if structural-context similarity. In *ACM SIGKDD*, 2002.

[16] A. Joshi, R. Kumar, B. Reed, and A. Tomkins. Anchor-based proximity measures. In *WWW '07*.

[17] G. Karypis and V. Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM J. Sci. Comput.*, 20(1):359–392, 1998.

[18] D. Liben-Nowell and J. Kleinberg. The link prediction problem for social networks. In *CIKM '03*, 2003.

[19] Q. Mei, D. Zhou, and K. Church. Query suggestion using hitting time. In *CIKM '08*.

[20] H. Qiu and E. R. Hancock. Image segmentation using commute times. In *Proc. BMVC*, 2005.

[21] H. Qiu and E. R. Hancock. Robust multi-body motion tracking using commute time clustering. In *ECCV'06*, 2006.

[22] P. Sarkar and A. Moore. Fast nearest-neighbor search in disk-resident graphs. Technical report:10-005, Machine Learning Department, Carnegie Mellon University, 2010.

[23] T. Sarlós, A. A. Benczúr, K. Csalogány, D. Fogaras, and B. Rácz. To randomize or not to randomize: space optimal summaries for hyperlink analysis. In *WWW*, 2006.

[24] A. D. Sarma, S. Gollapudi, and R. Panigrahy. Estimating pagerank on graph streams. In *PODS*, 2008.

[25] D. Spielman and S. Teng. Nearly-linear time algorithms for graph partitioning, graph sparsification, and solving linear systems. In *Proceedings of the STOC'04*.

[26] X. Zhu, Z. Ghahramani, and J. Lafferty. Semi-supervised learning using gaussian fields and harmonic functions. In *ICML, volume 20*, 2003.

7. APPENDIX

Symmetry of degree normalized PPV in undirected graphs. This follows directly from the reversibility of random walks.

$$v_i(j) = \alpha \sum_{t=0}^{\infty} (1-\alpha)^t P^t(i,j) = \frac{d_j}{d_i} \alpha \sum_{t=0}^{\infty} (1-\alpha)^t P^t(j,i)$$

$$\Rightarrow v_i(j)/d_j = v_j(i)/d_i \qquad (6)$$

Proof of theorem 3.1. Personalized pagerank of a start distribution r can be written as

$$PPV(r) = \alpha r + \alpha \sum_{t=1}^{\infty} (1-\alpha)^t (P^T)^t r$$

$$= \alpha r + (1-\alpha)PPV(P^T r) \qquad (7)$$

By turning node s into a sink, we are *only* changing the s^{th} row of P. We denote by r_s the indicator vector for node s. For $r = r_s$ we have personalized pagerank from node s. Essentially we are subtracting the entire row $P(s,:) = P^T r_s$ and adding back r_s. This is equivalent to subtracting the matrix vu^T from P, where v is r_s and u is defined as $P^T r_s - r_s$. Thus we have:

$$PPV(r) = \alpha(I - (1-\alpha)P^T + (1-\alpha)uv^T)^{-1}r$$

Let $M = I - (1-\alpha)P^T$. Hence $PPV(r) = \alpha M^{-1}r$. A straightforward application of the Sherman Morrison lemma gives

$$\widehat{PPV}(r) = PPV(r) - \alpha(1-\alpha)\frac{M^{-1}uv^T M^{-1}}{1 + (1-\alpha)v^T M^{-1}u}r$$

Note that: $M^{-1}u = 1/\alpha[PPV(P^T r_s) - PPV(r_s)]$. and also, $M^{-1}r = 1/\alpha PPV(r)$. We also have, $v^T PPV(r) = PPV(r,s)$.

Combining these facts with eq. (7) yields the following:

$$\widehat{PPV}(r) = PPV(r) - [PPV(r_s) - r_s]\frac{PPV(r,s)}{PPV(r_s,s)}$$

This leads to the element-wise error bound in theorem 3.1.

Proof of lemma 3.4. For proving the above we use a series of sink node operations on a graph and upper bound each term in the sum. For $j \leq k$ $S[j]$ denote the subset $\{s_1, ..., s_j\}$ of S. Also let $G \setminus S[j]$ denote a graph where we have made each of the nodes in $S[j]$ a sink. $S[0]$ is the empty set and $G \setminus S[0] = G$. Since we do *not* change the outgoing neighbors of any node when make a node into a sink, we have $G \setminus S[j] = (G \setminus S[j-1]) \setminus s_j$. Which leads to:

$$PPV^{G \setminus S[k-1]}(r,i) - PPV^{G \setminus S[k]}(r,i)$$
$$\leq \frac{PPV^{G \setminus S[k-1]}(s_k,i)PPV^{G \setminus S[k-1]}(r,s_k)}{PPV(s_k,s_k)}$$
$$\leq PPV^{G \setminus S[k-1]}(s_k,i)$$

The last step can be obtained by combining linearity of personalized pagerank with lemma 3.2. Now, using a telescoping sum:

$$PPV^G(r,i) - PPV^{G \setminus S[k]}(r,i) \leq \sum_{j=1}^{k} PPV^{G \setminus S[j-1]}(s_j,i)$$

The above equation also shows that by making a number of nodes sink the personalized pagerank value w.r.t any start distribution at a node can only decrease, which intuitively makes sense. Thus each term $PPV^{G \setminus S[k-1]}(s_k,i)$ can be upper bounded by $PPV^G(s_k,i)$, and $PPV^{G \setminus S[k-1]}(r,s_k)$ by $PPV^G(r,s_k)$. This gives us the following sum, which can be simplified using (6) as,

$$\sum_{j=1}^{k} PPV(s_j,i) = \sum_{j=1}^{k} \frac{d_i PPV(i,s_j)}{d(s_j)} \leq \frac{d_i}{\min_{s \in S} d_s}$$

The last step follows from the fact that $PPV(i,s_j)$, summed over j has to be smaller than one. This leads to the final result in lemma 3.4.

Proof of lemma 3.9. The error at any node i is given by $1/\alpha PPV_\alpha(\bar{\epsilon},i)$. From the definition of PPV with start distribution $\bar{\epsilon}$, we know,

$$PPV_\alpha(\bar{\epsilon},i) = \sum_j \sum_{t=0}^{\infty} \alpha(1-\alpha)^t (P^T)^t(i,j)\bar{\epsilon}(j)$$
$$\leq |\bar{\epsilon}|_\infty \sum_{t=0}^{\infty} \alpha(1-\alpha)^t \sum_j P^t(j,i)$$
$$= \sum_{t=0}^{\infty} \alpha(1-\alpha)^t \sum_j \frac{d_i P^t(i,j)}{d_j}$$
$$\leq \frac{d_i}{\delta} |\bar{\epsilon}|_\infty \sum_{t=0}^{\infty} \alpha(1-\alpha)^t$$
$$\leq \frac{d_i}{\delta} |\bar{\epsilon}|_\infty$$

The third step uses reversibility of random walks, i.e. $d_i P^t(i,j) = d_j P^t(j,i)$. δ is the minimum weighted degree. Since the maximum entry of $\bar{\epsilon}$ is $\frac{\alpha\epsilon}{1 - \sqrt{1-\alpha}}$, the result follows.